

Reactive Programming


J. Heinzlreiter
Version 2.3

Content

- Part 1: Introduction
 - Motivation for Reactive Programming
 - Traditional Concurrent Programming Techniques
 - Foundations of Functional Programming
- Part 2: Futures and Promises
- Part 3: Actors
 - Foundations of Actor Systems and Akka
 - Designing Actor Systems
 - Actor Supervision
 - Distributed Programming with Actors
- Part 4: Reactive Streams (Akka Streams)




Literature (1)

- [Odersky, 2021]: Martin Odersky, Lex Spoon, Bill Venners. **Programming in Scala**. Artima, 5th ed., 2021.
- [Pilquist et al., 2023]: Michael Pilquist, Rúnar Bjarnason, and Paul Chiusano. **Functional Programming in Scala**. Manning, 2nd ed., 2023.
- [Prokopec, 2017]: Aleksandar Prokopec. **Learning Concurrent Programming in Scala**. Packt Publishing, 2nd ed., 2017.



Literature (2)

- [Odersky et al., 2015]: Martin Odersky, Erik Meijer, Roland Khun: [Principles of Reactive Programming](#). Online Course, 2015 (see YouTube).
- [Roostenburg et al., 2017]: R. Roostenburg, R. Bakker, R. Williams. **Akka in Action**. Manning, 2017.
- [Lopez-Sancho Abraham, 2023]: F. Lopez-Sancho Abraham. **Akka in Action**. Manning, 2nd ed., 2023.
- [Kuhn, 2017]: R. Kuhn. **Reactive Design Patterns**. Manning, February 2017.




Motivation for Reactive Programming

[Odersky et al., 2015; What is Reactive Programming?]
[<http://www.reactivemanifesto.org>]

Moore's Law

- **Moore's law** is the observation that **the number of transistors** on integrated circuits **doubles** approximately every two years.
- Clock speed doesn't increase correspondingly anymore.
- Consequences:
 - Shift from single- to multi-core computers.
 - Concurrent programming is becoming more and more important.
 - Modern programming languages are becoming functional.





<https://github.com/karlrupp/microprocessor-trend-data>

"The Free Lunch Is Over"

- Article from Herb Sutter (2005): The Free Lunch Is Over. A Fundamental Turn Toward Concurrency in Software
 - *Instead of driving clock speeds and straight-line instruction throughput ever higher, they are instead turning en masse to **hyper-threading and multicore architectures**.*
 - ***Applications will increasingly need to be concurrent** if they want to fully exploit continuing exponential CPU throughput gains.*
 - *Efficiency and **performance optimization will get more, not less, important**.*
 - *The **vast majority of programmers today don't grok concurrency**, just as the vast majority of programmers 15 years ago didn't yet grok objects.*

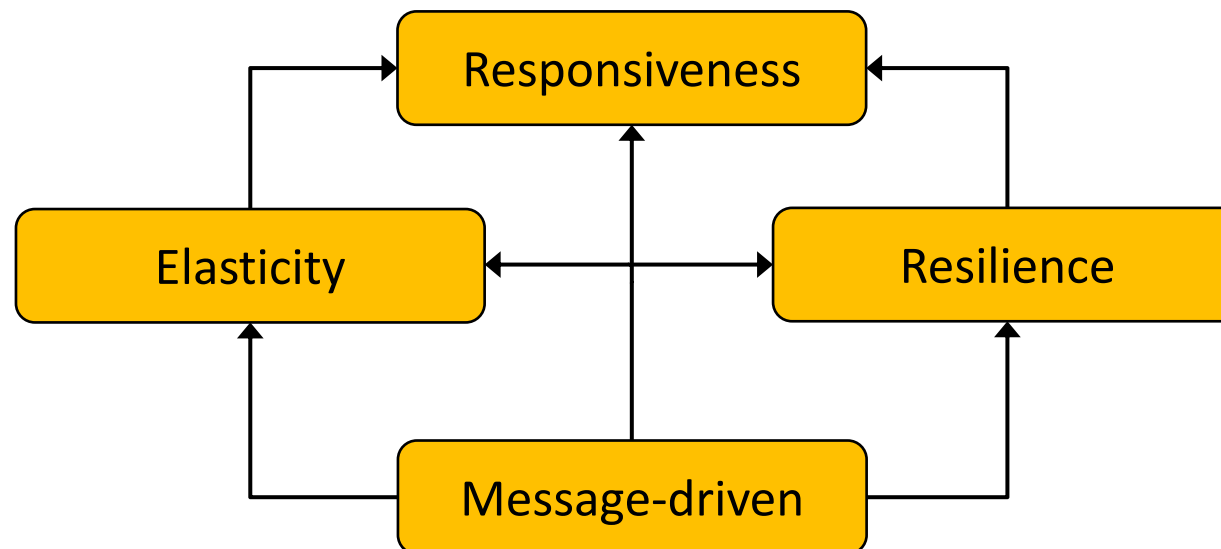
Concurrent vs. Parallel

- **Parallelism:** Techniques to accelerate programs by performing computations simultaneously.
 - Requires multiple CPUs
 - Computation units must be independent
 - Goal: Increase runtime efficiency
- **Concurrency:** Program makes progress on multiple tasks at the same time.
 - Concurrent programs can be executed on a single CPU, but can benefit from multiple CPUs
 - Goal: Allow efficient interaction with multiple external agents
 - Concurrency includes parallelism



Reactive Systems

- Definition of *reactive*: “Showing a response to a stimulus. Pupils are reactive to light” [Oxford Dictionaries]
- The *Reactive Manifesto*: Characteristics of Reactive Systems




Characteristics of Reactive Systems (1)

- **Message-driven:** react to events
 - Reactive Systems rely on asynchronous message-passing → no blocking.
 - Ensures loose coupling of components → location transparency.
 - Message queues enable load management.
 - Also, failures are communicated via messages.
- **Elasticity:** react to load
 - The system stays responsive under varying workload.
 - *Scale up:* make use of multiple cores, increase memory
 - *Scale out:* make use of multiple server nodes
 - Requirements:
 - Minimize mutable state
 - Location transparency, resilience

Characteristics of Reactive Systems (2)

- **Resilience:** react to failures
 - The system stays responsive in the face of failure.
 - Parts of the system can fail and recover without compromising the system as a whole.
 - Recovery of each component is delegated to another component.
 - Prerequisites:
 - Encapsulation of state
 - Supervisor hierarchies
- **Responsiveness:** react to users
 - System must provide real-time interaction with agents under load and in the presence of failures.
 - Prerequisites: Message-driven architecture, elasticity, resilience.



Traditional Concurrent Programming Techniques


[Prokopec, 2017; p. 27 ff.]

Overview

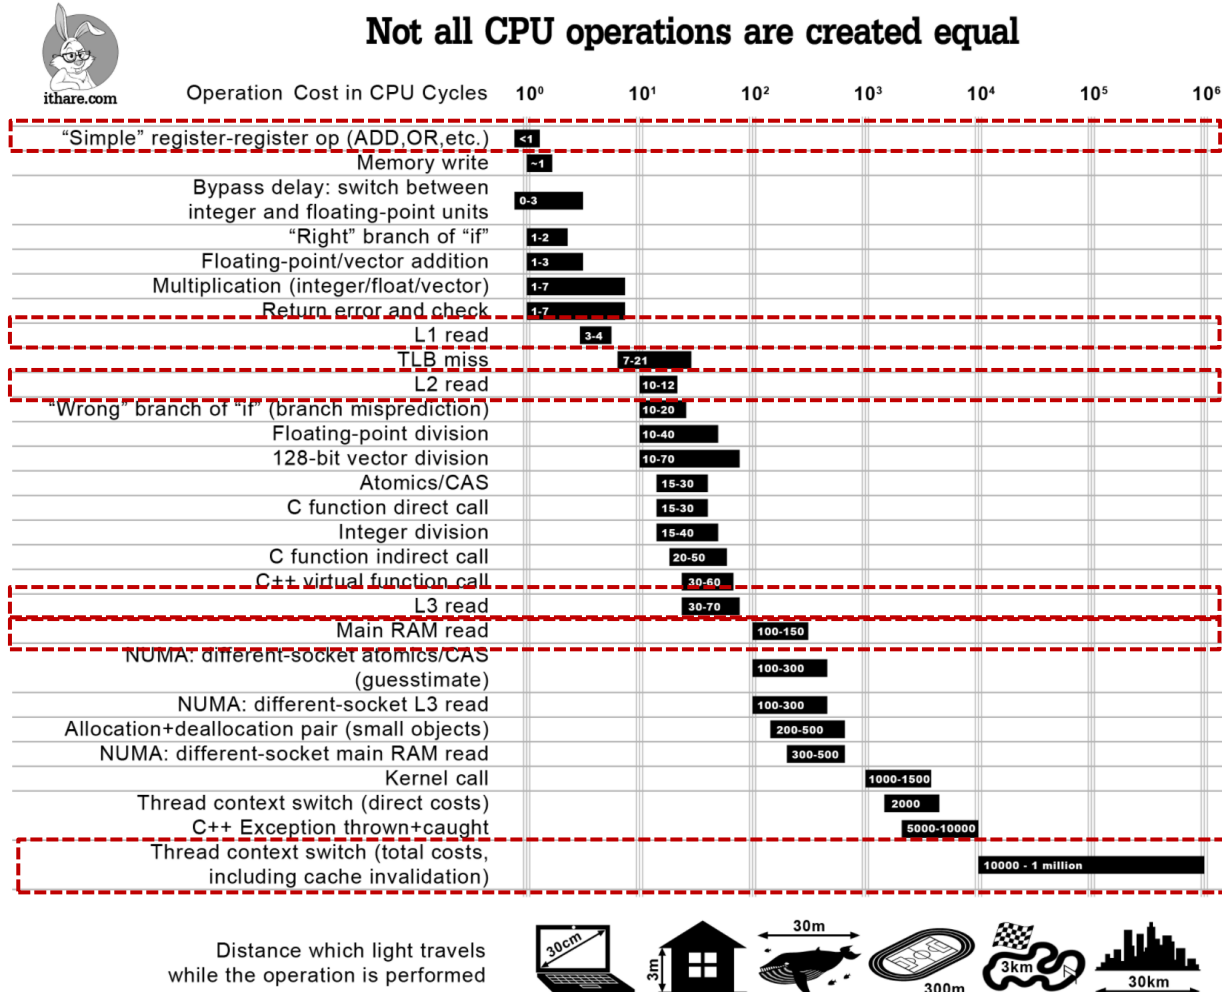
- Traditional low-level constructs for concurrent executions:
 - Processes and
 - Threads
- Means for communication:
 - Shared memory communication → synchronization required
 - Message passing → distributed systems
- Problem: low level and error prone
 - Deadlocks
 - Race conditions
 - Data races
 - Starvation
- Basic knowledge of low-level approaches is essential to understand high-level concepts.

Processes and Threads

- A *process* executes the instructions of a program.
 - Shares CPU and other resources with other processes.
 - Has its own share of memory.
- A process' instructions are executed concurrently in multiple *threads*.
 - Threads share the memory of the process they belong to.
 - Threads communicate by writing and reading to memory.
- Threads can be executed on different processors (cores).
 - Processors do not write directly to main memory.
 - They use *caches* to improve read/write performance.
 - One processor cannot access the cache of another processor.
- Java threads are directly mapped to *OS threads*.
- Scala inherits the thread model of Java.



Costs of different CPU operations



<http://ithare.com/infographics-operation-costs-in-cpu-clock-cycles>

Java Memory Model (JMM)

- Scala inherits the memory model from the JVM.
- The memory model defines *when writes to a variable are visible to other threads*.
- The compiler and the runtime perform various optimizations to gain performance.
 - *Registers* may be used as intermediate storage.
 - Data may be written to hierarchies of *caches*.
 - *Bytecode* statements may be *reordered*.
- The **rules** of the JMM define how threads interact through memory:
 - *Program order*: Program optimizations by the compiler must not alter the serial semantics within a thread.
 - *Locking*: Locks for the same synchronization variable must not overlap.
 - *Volatile fields*: A write to a volatile field is immediately visible to all threads.
 - *Thread start*: All actions in a thread are done after a call to `start()`.
 - Many other rules.

Working with Threads in Scala

- Thread creation

```
class MyThread extends Thread:  
  override def run(): Unit =  
    println("Executed in new thread.")  
  
val t = new MyThread  
t.start()
```

- Convenience method to execute code in separate thread

```
def doInThread(body: => Unit): Thread =  
  val t = new Thread:  
    override def run() = body  
  
  t.start()  
  t
```

```
val t = doInThread { println("Executed in new thread.") }
```

Execution of Multi-Threaded Code

- If multiple threads access the same memory locations, the program behavior is not deterministic anymore.

```
var currId = 0L
def generateUniqueId() =
  val newId = currId + 1 // (1)
  currId = newId         // (2)
  newId

val t1 = doInThread { println(generateUniqueId()) }
val t2 = doInThread { println(generateUniqueId()) }
// Output: 1 2 or 2 1 or 1 1
```

- The program works correctly if execution of statements (1) and (2) is not interfered by the other thread.
- t1 and t2 generate the same id if both threads concurrently read the field *currId*.
- We say that there is a *race condition* in a program when the output of a program depends on the execution schedule of the instructions.

Synchronization

- Every Java/Scala object has a special property called an *intrinsic lock or monitor*.
- Only one thread can *acquire* ownership of a lock.
- If a thread tries to acquire the lock while another thread owns the lock, this thread is *blocked*.
- If a thread executes the statement `x.synchronized { statements }`
 - it tries to acquire the lock for `x` which is released at the end of the block.
 - The thread is blocked while `x` is locked by another thread.
- The race condition in the example can be removed by locking the region accessed by multiple threads:

```
var currId = 0L
def generateUniqueId() = this.synchronized
  val newId = currId + 1
  currId = newId
  newId
```

Data Races and Reordering

- How can this program result in `x==1 && y==1`?

```
var a = false; var b = false
var x = -1;    var y = -1

val t1 = doInThread { a = true; y = if (b) 0 else 1 }
val t2 = doInThread { b = true; x = if (a) 0 else 1 }

t1.join(); t2.join()
assert(!(x == 1 && y == 1))
```

- It's required that `a` as well as `b` are `false` to produce `x==1 && y==1` as a result.
- Possible reasons for this “abnormal” behavior:
 - Reordering of statements*: This reordering has no impact on serial semantics but influences behavior of concurrent execution.
 - Data race*: Write of one thread (to cache) is not seen by the other thread.
- The `synchronized` keyword also guarantees that writes to memory of one thread are visible to all other threads.

Deadlocks

- Synchronization can result in deadlocks:

```
class Account(val name: String, var balance: Double);  
def transfer(amount: Double, a: Account, b: Account): Unit =  
  a.synchronized {  
    b.synchronized { a.balance -= amount; b.balance += amount }  
  }
```

```
var t1 = doInThread { for (i <- 0 until 1000)  
                      transfer(10, account1, account2) }  
var t2 = doInThread { for (i <- 0 until 1000)  
                      transfer(10, account2, account1) }
```

- Deadlock can be avoided when resources are always locked in the same order:

```
def transfer(amount: Double, a1: Account, a2: Account): Unit =  
  if (a1.name < a2.name)  
    a1.synchronized { a2.synchronized { ... } }  
  else  
    a2.synchronized { a1.synchronized { ... } }
```

- In practice it's difficult to ensure a consistent ordering of resources.

Guarded Blocks

- Busy-waiting can be avoided by using Java's *wait/notify* mechanism.

```
class Consumer(val tasks: Queue[() => Unit]) extends Thread:
  def addTask(task: => Unit) = tasks.synchronized {
    tasks.enqueue(() => task)
    tasks.notifyAll();
  }

  private def getTask(): () => Unit = tasks.synchronized {
    while (tasks.isEmpty) tasks.wait() // guarded block
    tasks.dequeue();
  }

  override def run() =
    while true do
      val task = getTask()
      task()
```

- Wait can cause *spurious wakeups*: Condition thread is waiting for is not met.
- Condition must be checked repetitively → *guarded block*.

Volatile Variables

- Reads and writes to volatile variables are atomic.
- Operations (eg. incrementation) are not atomic.
- Writes to volatile variables are immediately visible to all threads.
 - Variables are not cached or held in registers.
 - Compiler will not reorder instructions.
- Thread communication via volatile variables is very fast.
- Most common application: status flags


```
class Worker extends Thread:  
  @volatile  
  private var stopped = false  
  override def run() =  
    while !stopped do { /* do some work */ }  
  def shutdown() = stopped = true  
}
```

- This program may not terminate if `@volatile` is omitted.

Thread Pools

- Creating processes is very expensive, creating threads is expensive.
- Thread pools manage a set of threads.


ThreadPoolExecutor



Advantages of *ForkJoinPool*

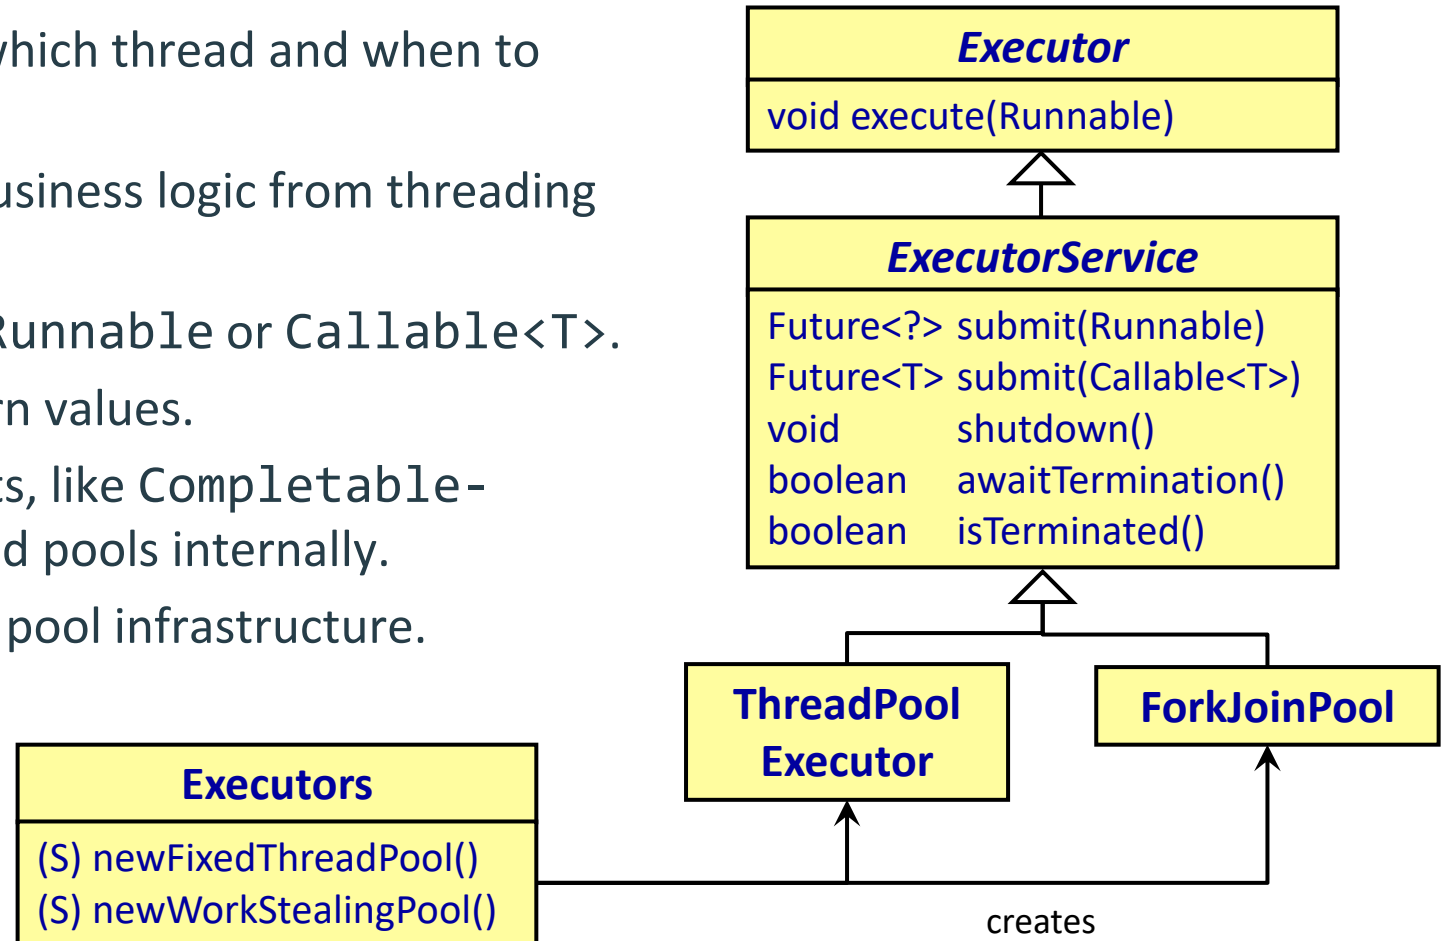
- Efficient for recursive tasks
- Less contention in local queue
- Load balancing achieved with *work stealing*

ForkJoinPool



Executor and ExecutorService

- Executor decides on which thread and when to execute a task.
- Executor decouples business logic from threading infrastructure.
- Tasks must implement `Runnable` or `Callable<T>`.
- `Callable` objects return values.
- More advanced concepts, like `CompletableFuture` leverage thread pools internally.
- Scala uses Java's thread pool infrastructure.



Using Executors in Scala

- Executor implementations can be used directly in Scala.

```
val executor = new java.util.concurrent.ForkJoinPool
executor.execute(
  () => println("This task is run asynchronously."))
```

- ExecutionContext is similar, but more Scala specific.
- Internally ExecutionContext uses a ForkJoinPool instance.

```
val execCtx = scala.concurrent.ExecutionContext.global
execCtx.execute(
  () => println("This task is run asynchronously."))
```

- Often an ExecutionContext object is passed implicitly to methods.
- A customized ForkJoinPool instance can also be used:

```
val execCtx = scala.concurrent.ExecutionContext.fromExecutorService(
  new java.util.concurrent.ForkJoinPool(2))
```

Atomic Variables

- Atomic variables support complex operations (more than one read/write), that are executed atomically.


```
private val uid = new AtomicLong(0L)
def getId(): Long = uid.incrementAndGet()
```

- Operations are implemented *lock-free*
 - Operations are very fast
 - No danger of deadlocks
- Operations are implemented in terms of a fundamental atomic operation: *compareAndSet* (also called *compare-and-swap*).
 - Takes the expected previous value and the new value
 - Fails if previous value changed → operation has to be repeated


```
private val currId = new AtomicLong(0L)
@tailrec def generateUniqueId(): Long =
  val oldId = currId.get
  val newId = oldId + 1
  if currId.compareAndSet(oldId, newId) then newId else generateUniqueId()
```


Concurrent Collections

- *Producer-Consumer* is a common pattern in concurrent programming.
 - Work items need to be buffered in *concurrent queue*



- Java's concurrent queues can also be used in Scala.
- `BlockingQueue` supports three types of operations, depending on the behavior when queue is empty or full:
 - Operations throw an exception: `add/remove/element`
 - Operations return a special value: `offer/poll/peek`
 - Operations block: `put/take`
- Iterators are *weakly consistent*: modifications might or might not be reflected in iterator.





Foundations of Functional Programming

[Pilquist et al., 2023]

High-Order Functions

- Functions which take other functions as parameters are called *high-order functions*.
- It's easy to create and to handle functions: Functions are first-class citizens in Scala.
- Example:

```
def sumSquares(a: Int, b: Int): Int =  
  if a > b then 0 else a * a + sumSquares(a + 1, b)  
def sumPowerOfTwo(a: Int, b: Int): Int =  
  if a > b then 0 else Math.pow(2, a).toInt + sumPowerOfTwo(a + 1, b)
```

```
def sum(f: Int => Int, a: Int, b: Int): Int =  
  if (a > b) 0 else f(a) + sum(f, a + 1, b)  
def sumSquares(a: Int, b: Int): Int = sum(x => x * x, a, b)  
def sumPowerOfTwo(a: Int, b: Int): Int = sum(Math.pow(2, _).toInt, a, b)
```

- Parameterization of functions by other functions is a typical design pattern in functional programming.

Immutable State

- Typical procedural Java code:

```
public class Product { ... }  
public static List<String> outOfStock(List<Product> allProds) {  
    List<Product> prods = new ArrayList<>();  
    for (Product p : allProds) if (p.inStock < MIN_STOCK) prods.add(p);  
    Collections.sort(prods, new Comparator<Product>() { ... });  
    List<String> prodNames = new ArrayList<String>();  
    for (String p : prods) prodNames.add(p.getName());  
    return prodNames;  
}
```

- Functional Scala code is cleaner, because
 - no state has to be tracked,
 - code describes *what* has to be done and *not how* it has to be done.

```
case class Product(id: Int, name: String, inStock: Int)  
def outOfStock(prods: Seq[Product]) = prods.filter(_.inStock < MIN_STOCK)  
                                           .sortBy(_.id)  
                                           .map(_.name)
```

Applying Functions to Collections

- `filter` composes a collection of all elements for which a given predicate holds:

```
val evenNumbers = List(1, 2, 3, 4, 5) filter (_ % 2 == 0) // List(2, 4)
```

- `map` applies a function to each element of a collection:

```
val lowerCase = List("a", "b", "c")  
val upperCase = lowerCase map (_.toUpperCase) // List("A", "B", "C")
```

```
val lowerUpper = lowerCase map (str => List(str, str.toUpperCase))  
// List(List("a", "A"), (List("b", "B"), List("c", "C"))
```

- `flatMap`
 - Applies a function that generates a sequence out of each element of a collection.
 - Flattens the resulting sequence of sequences to a (flat) sequence.

```
val lowerUpper = lowerCase flatMap (str => List(str, str.toUpperCase))  
// List("a", "A", "b", "B", "c", "C")
```


For Expressions (1)

- map and flatMap can be cascaded → code may be hard to understand:

```
case class Person(name: String, isMale: Boolean, children: Seq[Person])  
val persons: Seq[Person] = List(adam, eve, cain, abel)  
val parentChildPairs = persons flatMap (p =>  
    p.children map (c => (p.name, c.name)))
```

- yields in List((Adam,Cain), (Eve,Cain), (Eve,Abel))
 - using map would yield in List(List((Adam,Cain)), ..., List())
- The following for expression does exactly the same:

```
val parentChildPairs =  
    for (p <- persons;  
         c <- p.children)  
    yield (p.name, c.name)
```

- The compiler translates the second query into the first one.
- The second query is much clearer to read.

For Expressions (2)

- Filters

```
val motherChildPairs = persons.filter (p => !p.isMale)
                                flatMap (p =>
                                    p.children map (c => (p.name, c.name)))
```

can also be mapped to for expressions

```
val motherChildPairs =
  for (p <- persons if (!p.isMale);
      c <- p.children) yield (p.name, c.name)
```

- For expressions with side effects

```
for (p <- persons) println(p.name)
```

are mapped to the function foreach

```
persons.foreach (p => println(p.name));
```

- A for expression can be applied to all domains that provide map, flatMap, filter/withFilter (and foreach).

Monads

- There are many types that define `map` and `flatMap`.
- A type `M` that supports the three operations `flatMap`, `map`, and `unit` is called a monad:

```
trait M[T]:  
  def flatMap[U](f: T => M[U]): M[U]  
  def map[U]    (f: T => U    ): M[U]  
  
def unit[T](x: T): M[T]
```

- `map` can be defined in terms of `flatMap` and `unit`:

```
def map[U](f: T => U) : M[U] = m flatMap (x => unit(f(x)))
```

- In a stricter definition monads must fulfill certain *laws*.
- In other functional languages `flatMap` is called `bind`.
- A monad need not be a collection. Often monads handle a *single value*.
- In Scala the name of the `unit` function is monad-specific.

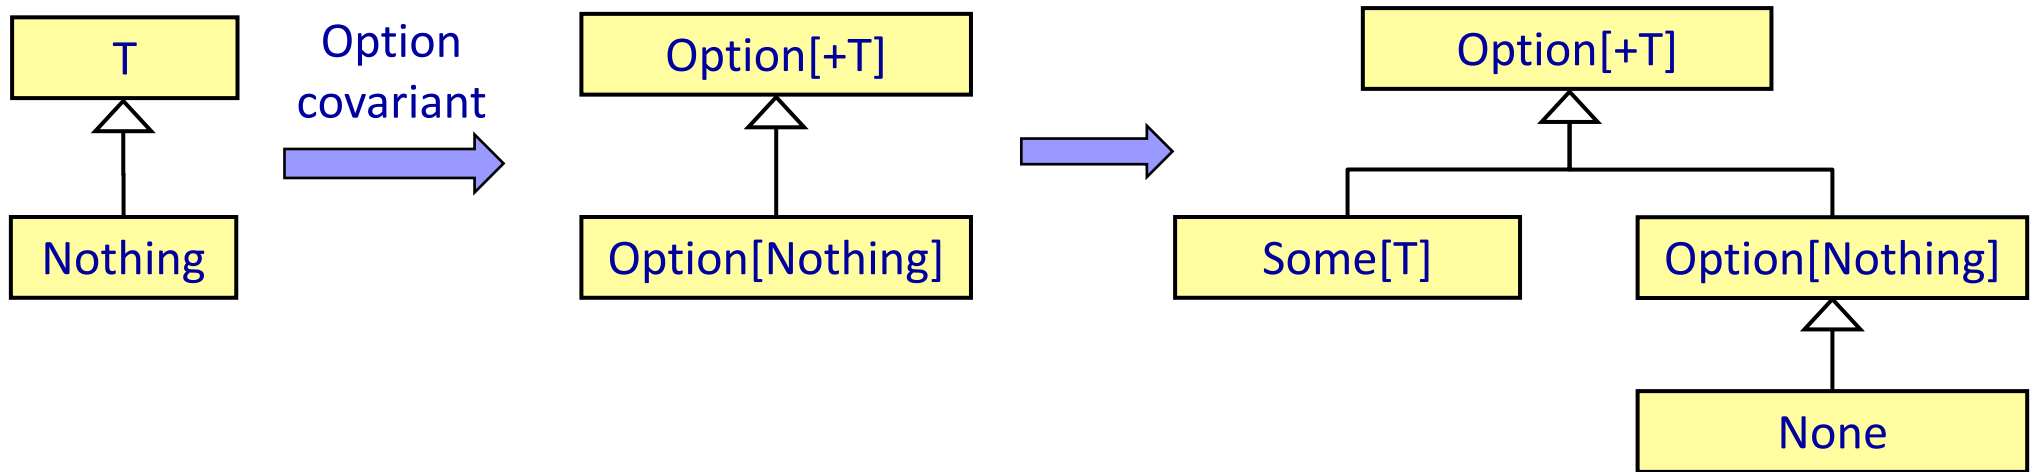
The Option Monad

- Options instances hold values of two forms:
 - Some(x) where x is the actual value
 - None represents an undefined value

```
abstract class Option[+T]:  
  def flatMap[U](f: T => Option[U]): Option[U] = this match  
    case Some(x) => f(x)  
    case None => None  
  
  def map[U](f: T => U): Option[U] = this match  
    case Some(x) => Some(f(x))  
    case None => None  
  
case class Some[+T](element : T) extends Option[T]  
case object None extends Option[Nothing]
```

- Advantages over null
 - Cleaner interfaces: It's obvious that value may be undefined.
 - Checking for undefined value is enforced.

The Option Monad: Embedding into the Type System



- Application:

```
val op1: Option[String] = Some("Hello")
val op2: Option[String] = None
```

The Option Monad: map

- The unit operations `Some` and `None` are generators for monad values:

```
val o1: Option[Double] = Some(9)
val o2: Option[Double] = None
```

- `map` transforms `Option` values to other `Option` values:

```
def toDouble(value: String): Double = value.toDouble
Some("9") map (n => toDouble(n)) // → Some(9.0)
None      map (n => toDouble(n)) // → None
```

- `map` allows for *function composition*:

```
def sqrt(value: Double): Double = math.sqrt(value)
Some("9") map (n => toDouble(n)) map (n => sqrt(n)) // → Some(3.0)
```

- Applying a function to a `None` value does not result in an exception:

```
None map (n => toDouble(n)) map (n => sqrt(n)) // → None
```

The Option Monad: flatMap

- map is not always useful when applied to option values because it generates “nested objects”:

```
def toDouble(value : String) : Option[Double] =  
  try { Some(value.toDouble) }  
  catch { case nfe : NumberFormatException => None }  
  
Some("9") map (n => toDouble(n)) // → Some(Some(9.0))  
Some("x") map (n => toDouble(n)) // → Some(None)
```

- This can be prevented by applying flatMap:

```
Some("9") flatMap (n => toDouble(n)) // → Some(9.0)  
Some("x") flatMap (n => toDouble(n)) // → None
```

- In this manner functions can also be composed:

```
def sqrt(value : Double) : Option[Double] = value match { ... }  
  
Some("9") flatMap (n => toDouble(n)) flatMap (n => sqrt(n)) // → Some(3.0)  
Some("x") flatMap (n => toDouble(n)) flatMap (n => sqrt(n)) // → None  
Some("-9") flatMap (n => toDouble(n)) flatMap (n => sqrt(n)) // → None
```

Monad Laws

■ Associativity

```
(m flatMap f) flatMap g == m flatMap (x => f(x) flatMap g)
```

```
(Some("4") flatMap (toDouble(_)) flatMap (sqrt(_)) ==  
  Some("4") flatMap (x => toDouble(x) flatMap (sqrt(_))) == Some(2.0)  
(None flatMap (toDouble(_)) flatMap (sqrt(_)) ==  
  None flatMap (x => toDouble(x) flatMap (sqrt(_))) == None
```

■ Left unit

```
unit(x) flatMap f == f(x)
```

```
Some("1") flatMap (toDouble(_)) == toDouble("1") == Some(1.0)
```

■ Right unit

```
m flatMap unit == m
```

```
Some("1") flatMap (Some(_)) == Some("1")  
None flatMap (Some(_)) == None
```


Monads and For Expressions

- Because monads support `map` and `flatMap` they can be used in for expressions.
- Function composition expressions like

```
val res = Some("4") flatMap (toDouble(_))  
                      flatMap (sqrt(_))
```

can also be written as for expressions:

```
val res =  
  for (o1 <- Some("4");  
       o2 <- toDouble(o1);  
       o3 <- sqrt(o2)) yield o3;
```

- Advantage: better readability

The Try Monad (1)

- Try can have two types of values:
 - Success(x): Computation ended successfully with result x.
 - Failure(ex): Computation failed with Exception ex.

```
abstract class Try[+T]
case class Success[T](x: T) extends Try[T]
case class Failure(ex: Throwable) extends Try[Nothing]
object Try:
  def apply[T](expr: => T): Try[T] =
    try
      Success(expr)
    catch
      case NonFatal(ex) => Failure(ex)
```

- Fatal exceptions: OutOfMemoryError, StackOverflowError, InterruptedException, etc.
- Application:
 - Composition of functions that may fail
 - Pass results between functions, threads and processes

The Try Monad (2)

- Like each monad Try must define `flatMap` and `map`:

```
abstract class Try[+T]:  
  def flatMap[U](f: T => Try[U]): Try[U] = this match  
    case Success(x) => try f(x) catch { case NonFatal(ex) => Failure(ex) }  
    case fail: Failure => fail  
  
  def map[U](f: T => U): Try[U] = this match  
    case Success(x) => Try(f(x)) // Try.apply(f(x))  
    case fail: Failure => fail
```

- In the stricter interpretation Try is not a monad because the *left unit law does not hold*:

```
Try(expr) flatMap f != f(expr)
```

- The left-hand side never throws a (non-fatal) exception whereas the right-hand side will raise any exception thrown by `f`.

The Try Monad: map

- The unit operations `Success` and `Failure` are generators for monad values:

```
val success: Try[String] = Success("9")
val failure: Try[String] = Failure(IllegalArgumentException())
```

- `map` transforms `Try` values to other `Try` values:

```
def toDouble1(value: String): Double = value.toDouble
Success("9") map (n => toDouble1(n)) // → Success(9.0)
Failure(ex)  map (n => toDouble1(n)) // → Failure(ex)
```

- `map` allows for *function composition*:

```
def sqrt1(value: Double): Double = value match
  case v: Double if v >= 0 => math.sqrt(value)
  case _                  => throw IllegalArgumentException()
Success("9") map (n => toDouble1(n)) map (n => sqrt1(n)) // → Success(3.0)
```

- Applying a function to a `Failure` value does not result in an exception:

```
Failure(ex) map (n => toDouble1(n)) map (n => sqrt1(n)) // → Failure(ex)
```

The Try Monad: flatMap

- map is not always useful when applied to Try values because it generates “nested objects”:

```
def toDouble2(value: String): Try[Double] = Try { value.toDouble }  
Success("9") map (n => toDouble2(n)) // → Success(Success(9.0))  
Success("x") map (n => toDouble2(n)) // → Success(Failure(NumberFormatException))
```

- This can be prevented by applying flatMap:

```
Success("9") flatMap (n => toDouble2(n)) // → Success(9.0)  
Success("x") flatMap (n => toDouble2(n)) // → Failure(NumberFormatException)
```

- In this manner functions can also be composed:

```
def sqrt2(value: Double): Try[Double] = value match  
  case v : Double if v>=0 => Success(math.sqrt(value))  
  case _                  => Failure(IllegalArgumentException())  
Success("9") flatMap (n => toDouble2(n)) flatMap (n => sqrt2(n))  
                                     // → Success(3.0)  
Success("x") flatMap (n => toDouble2(n)) flatMap (n => sqrt2(n))  
                                     // → Failure(NumberFormatException)  
Success("-9") flatMap (n => toDouble2(n)) flatMap (n => sqrt2(n))  
                                     // → Failure(IllegalArgumentException)
```

The Try Monad: For Expressions

- Because Try supports map and flatMap, Try can be used in for expressions:

```
val r1 = for (d <- toDouble2("9")) yield sqrt1(d)
      // → Success(3.0)
val r2 = for (d <- toDouble2("9"); s <- sqrt2(d)) yield s
      // → Success(3.0)
val r3 = for (d <- ("x"); s <- sqrt2(d)) yield s
      // → Failure(NumberFormatException)
val r4 = for (d <- toDouble2(" 9"); s <- sqrt2(d)) yield s
      // → Failure(IllegalArgumentException)
```

- Try objects can be processed by pattern matching:

```
r1 match
  case Success(x) => println(s"Success($x)")
  case Failure(ex) => println(s"Failure($ex)")
```

Futures and Promises

[Prokopec, 2017; p. 101 ff.]

Asynchronous Programming

- In certain situations, threads have to wait for external resources (e. g. data arriving over the network).
- Traditional programs *block* threads in such cases.
- Possible problems with *blocking synchronization*:
 - Poor utilization of threads
 - Danger of deadlocks
 - Starvation of threads pools
- Asynchronous programming:
 - Thread has to wait for resource → separate computation is scheduled.
 - Thread continues.
 - Original computation is continued when resource is available.
 - Prevents expensive context switches.
- Futures and Promises are abstractions that support asynchronous programming.

Futures

- A future

```
trait Future[T]:  
  def value: Option[Try[T]]  
object Future:  
  def apply[T](computation: => T): Future[T]
```

- is a placeholder for a value of type T → *future value*
- executes an asynchronous computation → *future computation*

- Starting a future computation:

```
val f : Future[Int] = Future {  
  computeNthPrime(1000)  
} // shorthand for Future.apply(...)
```

- Accessing the future value (never do this in this way in production code):

```
println(f.value) // → None  
while (! f.isCompleted) { Thread.sleep(100) }  
println(f.value) // → Some(Success(7919))
```

ExecutionContext

- Future computations are carried out in an ExecutionContext.
- The ExecutionContext is passed as the 2nd parameter to Future.apply:

```
object Future:  
  def apply[T](computation :=>T)  
    (using executor: ExecutionContext): Future[T]
```

- The Scala library provides a default implementation for this implicit parameter:

```
import scala.concurrent.ExecutionContext.Implicits.global
```

- The global ExecutionContext is backed by ForkJoinPool
 - The actual/minimum/maximum number of threads is set to Runtime.availableProcessors.
- The ExecutionContext can be easily replaced:

```
given ExecutionContext =  
  ExecutionContext.fromExecutor(Executors.newFixedThreadPool(10))
```

Future callbacks (1)

- Callbacks are (partial) functions that are called when the future computation succeeds or fails.

```
trait Future[T]:  
  def foreach(f: T => U)           // success  
  def failed: Future[Throwable]    // Failure(ex) → Success(ex)  
  def onComplete(f: Try[T] => Unit) // success or failure
```

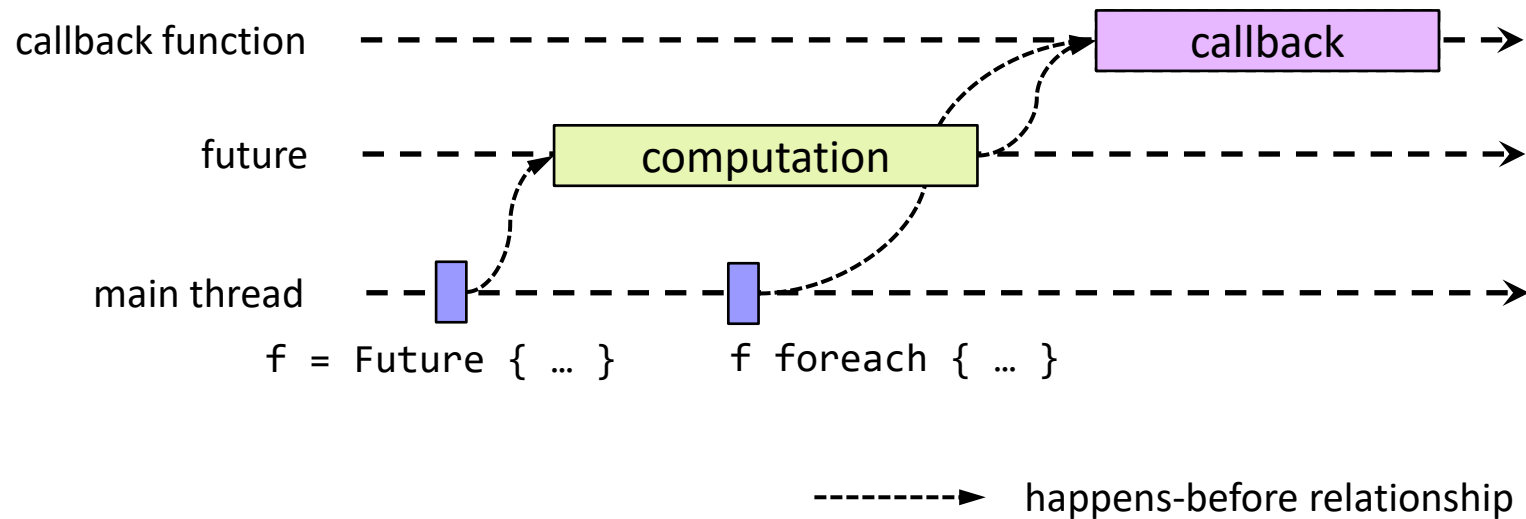
- Futures callbacks are called *eventually* after the computation is finished and independently of other callbacks.
- Example:

```
val f: Future[List[Int]] = Future { getPrimes(100) }  
f foreach { primes => println(primes.mkString(", ")) }  
f.failed foreach { ex => println(s"Failed with $ex") }  
f onComplete {  
  case Success(primes) => println(primes.mkString(", "))  
  case Failure(ex)      => println(s"Failed with $ex")  
}
```

Future callbacks (2)


- Installing a callback is a non-blocking operation (like future creation).
- Callback may or may not be executed in a separate thread.

```
val f = Future { computation }  
f foreach { callback }
```




Functional Composition of Futures (1)

- Problems related to callbacks
 - Cascading futures becomes cumbersome (nested callbacks)
 - Combining the results of multiple futures is difficult
 - Error handling becomes confusing
- Functional composition is a pattern in which functions are cascaded using high-order functions, so called *combinators*.
- Monads provide combinators and are therefore an important means for functional composition.



- Futures support combinators like `map`, `flatMap`, `filter`, etc. that can be used to build pipelines of future computations.



Functional Composition of Futures (2)

- Combinators generate futures out of futures → combinators don't block
- Combinators provided by Future:

```
trait Future[T]:  
  def map[S]      (f: T => S)          : Future[S]  
  def flatMap[S](f: T => Future[S]) : Future[S]  
  def filter      (p: T => Boolean)    : Future[T]  
  def recover[U >: T](pf: PartialFunction[Throwable, U]) : Future[U]
```

- map: Apply function f to future value. f is executed in a new future. Exceptions are propagated to the new future.
- flatMap: Similar to map, but result is not nested
 - val f1: Future[Int] = Future { 1 };
val f2 = f1 map { i => Future { i+1 } } → f2: Future[Future[Int]]
 - val f1: Future[Int] = Future { 1 };
val f2 = f1 flatMap { i => Future { i+1 } } → f2: Future[Int]
- recover: Function pf handles exceptions thrown in future computations. Valid results will be propagated to the new future.

Functional Composition of Futures (3)

- Example for recover:

```
val f1 = Future (6 / 0) recover { case e: ArithmeticException => 0 }  
    → Success(0)  
val f2 = Future (6 / 0) recover { case e: NullPointerException => 0 }  
    → Failure(ArithmeticException)  
val f3 = Future (6 / 2) recover { case e: ArithmeticException => 0 }  
    → Success(3)
```

- Example for map:

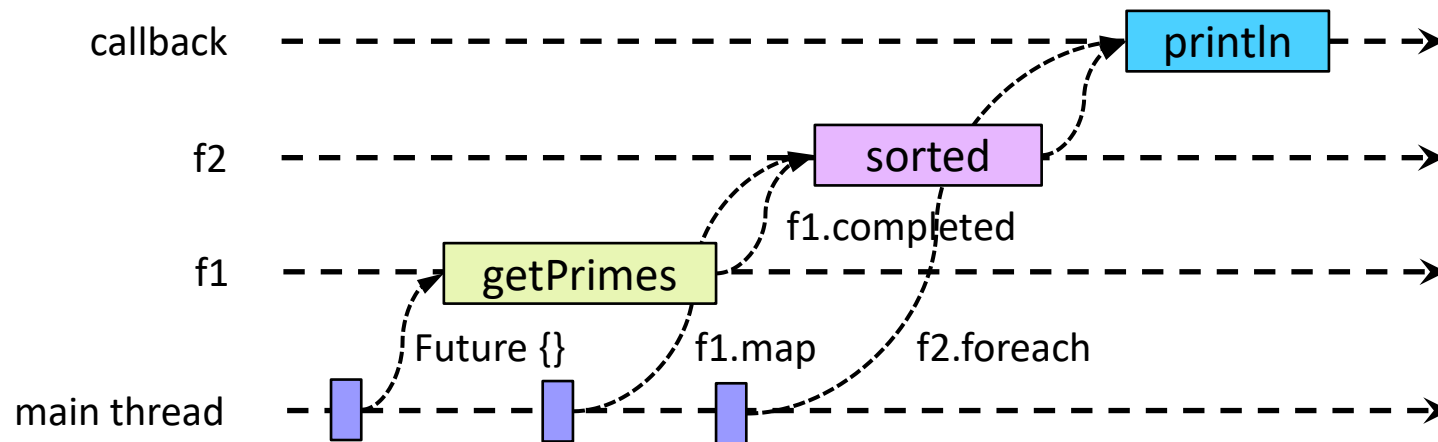
```
val f1: Future[List[Int]] = Future { getPrimes(100) }  
val f2: Future[List[Int]] = f1 map { primes => primes.sorted }  
f2 foreach { primes => println(primes.mkString(", ")) }
```

- Example for flatMap:

```
val f1 = Future { getPrimes(2, 50) }  
val f2 = Future { getPrimes(51, 100) }  
val f3 = f2 flatMap { primes2 => f1 map { primes1 => primes1 ++ primes2 } }  
val f4 = f3 map { primes => primes.sorted }  
f4 foreach { primes => println(primes.mkString(", ")) }
```

Semantics of the map Combinator

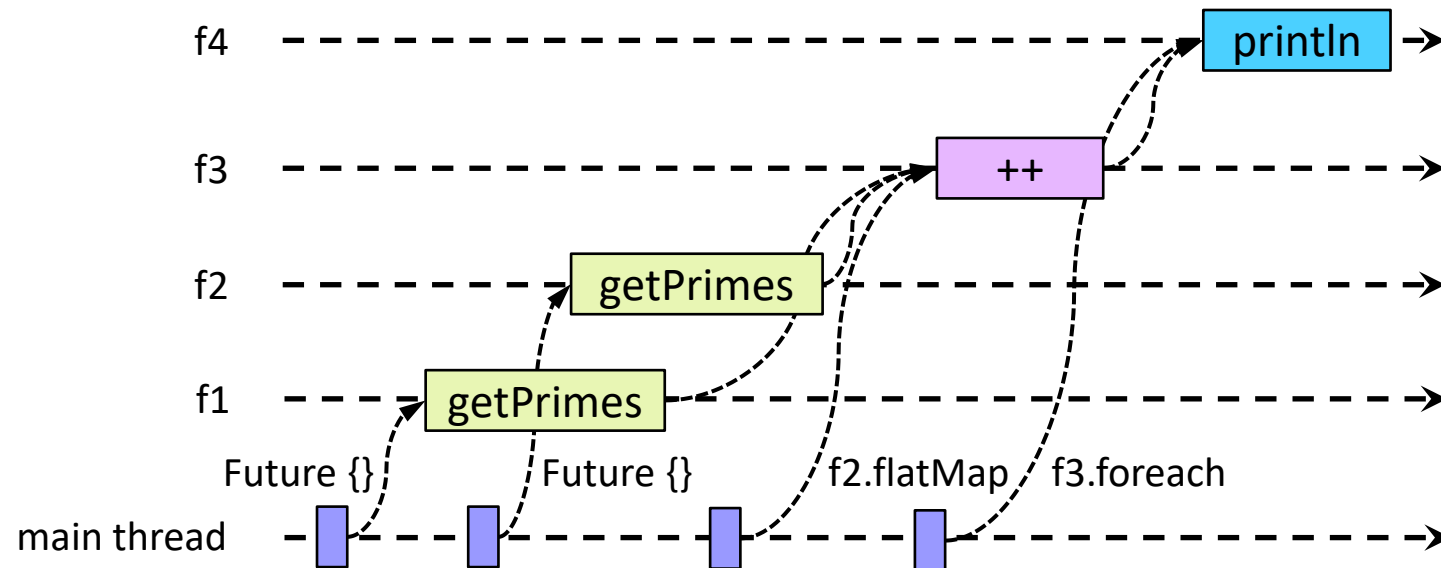
```
val f1: Future[List[Int]] = Future { getPrimes(100) }  
val f2: Future[List[Int]] = f1 map { primes => primes.sorted }  
f2 foreach { primes => println(primes.mkString(", ")) }
```



Futures and combinators may (but need not) be executed in separate threads.

Semantics of the flatMap Combinator (1)

```
val f1 = Future { getPrimes(2, 50) }  
val f2 = Future { getPrimes(51, 100) }  
val f3 = f2 flatMap { primes2 => f1 map { primes1 => primes1 ++ primes2 } }  
f3 foreach { primes => println(primes.mkString(", ")) }
```



Semantics of the flatMap Combinator (2)

- Difference between `Future.flatMap` and `Future.map`:
 - Let `f: Future[U]`, `g: U => V`
 - `val res = f flatMap { u => Future { g(u) } }`
 - Execution order:
 - Future `f` completes with a future value of type `U`.
 - Future `{ g(u) }` completes with a future value of type `V`.
 - Future `res` completes with a future value of type `V`.
 - Return type: `Future[V]`.
 - `val res = f map { u => Future { g(u) } }`
 - Execution order:
 - Future `f` completes,
 - Future `res` completes with a future value of type `Future[V]`,
 - Future `{ g(u) }` completes.
 - Return type: `Future[Future[V]]`.

Futures and For Expressions

- Since futures support the methods `map`, `flatMap`, `foreach` (and `withFilter`), futures can be combined in for expressions.
- Therefore, this program fragment

```
val f1 = Future { getPrimes(2, 50) }  
val f2 = Future { getPrimes(51, 100) }  
val f3 = f2 flatMap { primes2 => f1 map { primes1 => primes1 ++ primes2 } }  
f3 foreach { primes => println(primes.mkString(", ")) }
```

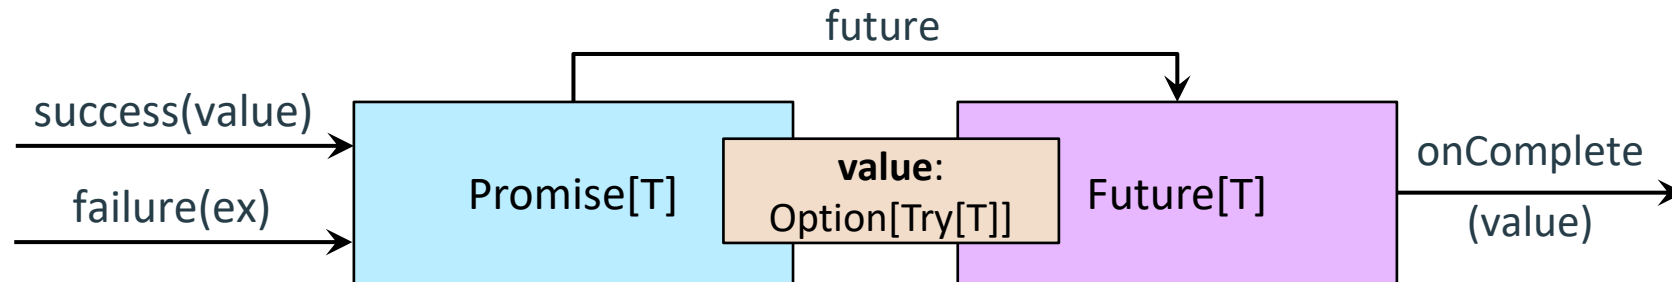
can be rewritten to

```
val f1 = Future { getPrimes(2, 50) }  
val f2 = Future { getPrimes(51, 100) }  
val f3 = for (primes1 <- f1;  
              primes2 <- f2) yield primes1 ++ primes2  
for (primes <- f3) println(primes.mkString(", "))
```

- Advantage: Program structure is easier to understand.

Promises (1)

- Futures and promises represent two aspects of a *single-assignment variable*:
 - The *promise* allows you to assign a value to the future.
 - The *future* allows you to read the value.



- Example usage of promises:

```
val p = Promise[String]()
p.future onComplete {
  case Success(v) => println(v)
  case Failure(ex) => println(ex)
}
p success "Hello"
// p.failure new Exception("failed")
```

Promises (2)

- A new value can only be *assigned once* to a future.
 - `success`, `failed`, and `complete` throw an exception if one tries to override a value.
 - `trySuccess`, `tryFailed`, and `tryComplete` return a value to indicate that the assignment was successful.
- It's easy to implement futures by means of promises:

```
import ExecutionContext.Implicits.global
def doInFuture[T](computation: => T): Future[T] =
  val p = Promise[T]()
  global.execute(() =>
    try
      val result = computation
      p success result
    catch
      case NonFatal(e) => p failure e
  )
  p.future
```

Future-Callback Bridge

- The control flow of callback-based APIs is not apparent from the code.
- The API dictates the control flow → *inversion of control*.
- Promises can be used to bridge the gap between callback-based APIs and futures.

```
case class Event(value: Any)
@FunctionalInterface
trait EventListener:
  def onEvent(evt: Event)
```

```
class EventSource:
  def addEventListener(l: EventListener) = ...
```

```
def eventOccurred(): Future[Event] =
  val p = Promise[Event]()
  val source = new EventSource
  source.addEventListener((evt: Event) => p trySuccess evt)
  p.future
```

```
eventOccurred() foreach { evt => println(s"event [$evt] occurred.") }
```

Futures and Blocking

- If all threads of the thread pool are blocked, no further future computation can be started.

```
var futures = for (_ <- 1 to noProcessors)
    yield Future { Thread.sleep(1000) }
val f = Future { ... }
```

- Computation of future `f` will only start after 1000 ms.
- This situation is called *thread starvation*: shared resources are made unavailable for long periods by greedy threads.
- Avoid blocking whenever possible.
- If blocking is absolutely necessary, enclose code in a `blocking` call.

```
var futures = for (_ <- 1 to noProcessors)
    yield Future { blocking { Thread.sleep(1000) } }
val f = Future { ... }
```

- The execution context will provide additional worker threads if necessary.

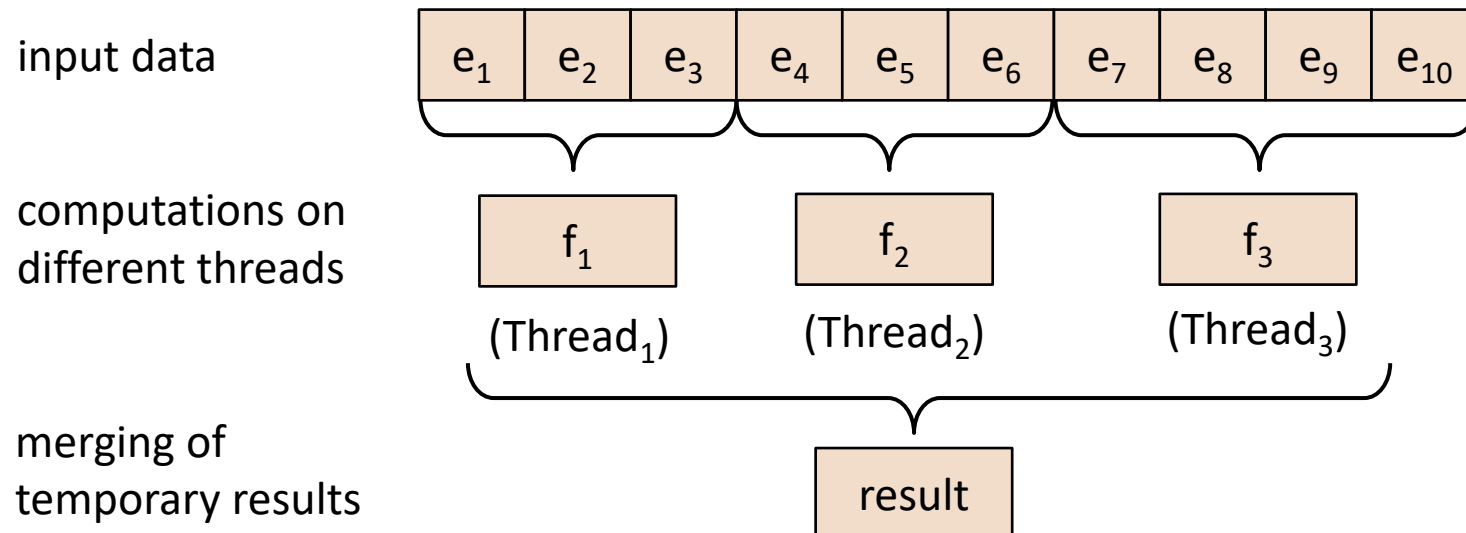


Data Parallel Collections

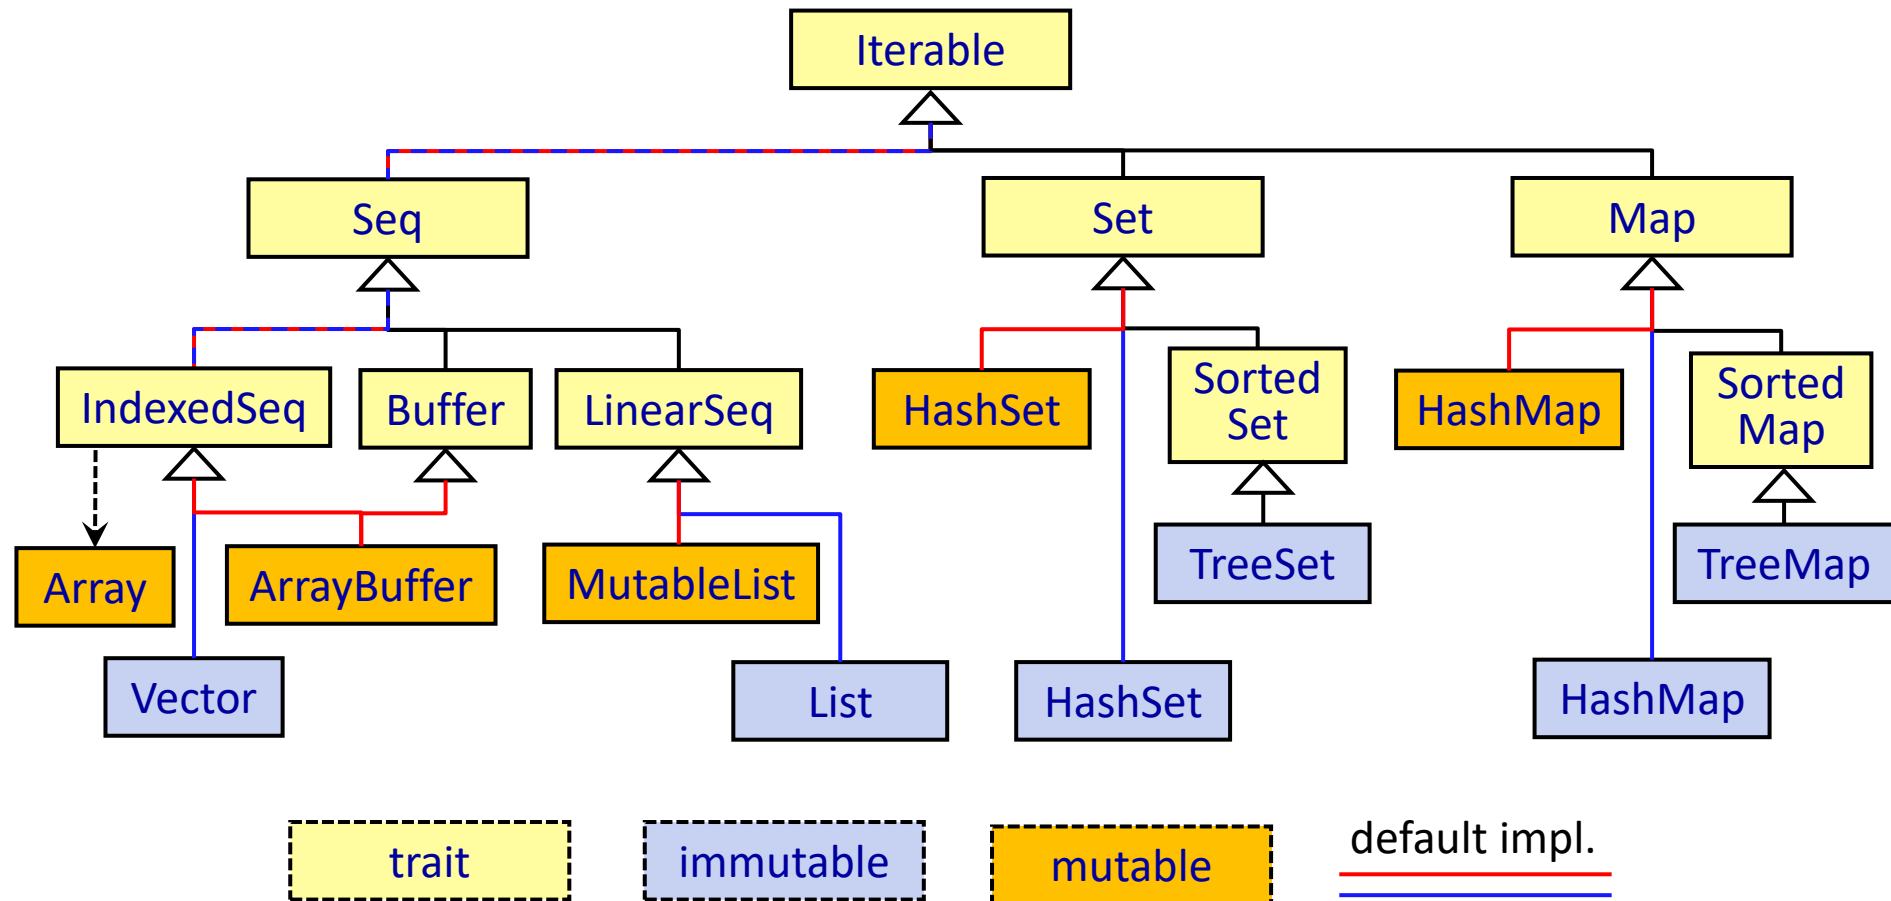
[Prokopec, 2017; p. 137 ff.]

Data Parallelism

- There are two forms of parallelism
 - **Task parallelism** (control parallelism) focusses on distributed tasks which are performed across different threads and processors.
 - **Data parallelism** focusses on performing the same operations on different data elements.
- Basic principle of data parallelism



Scala's Collection Hierarchy



Semantics of Scala Collections

- **Iterable:**

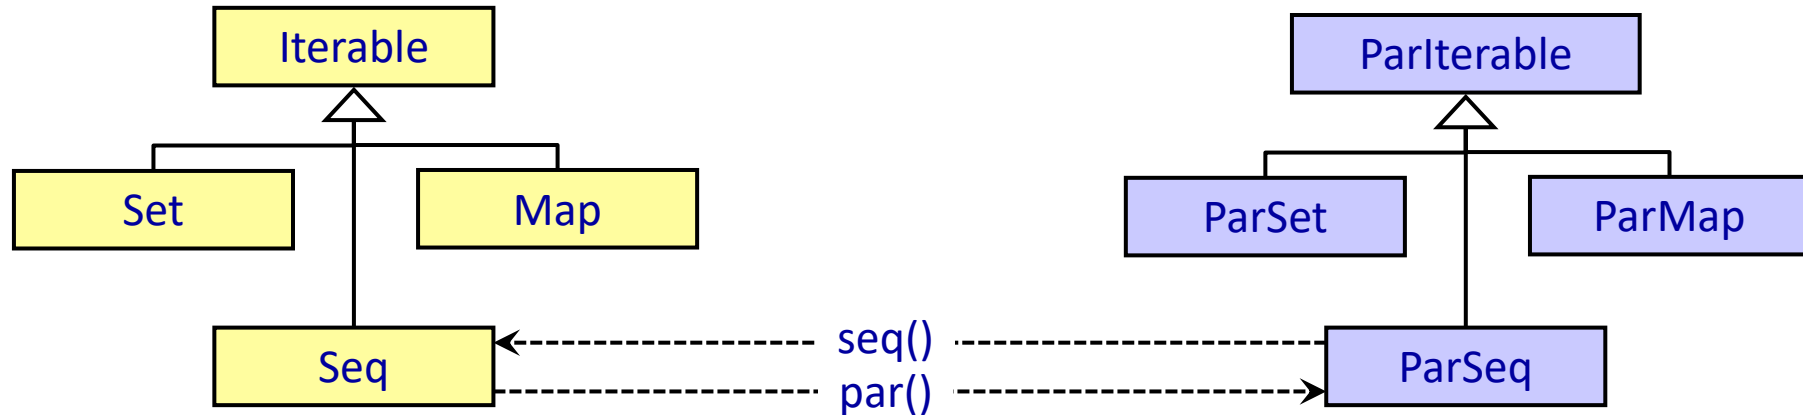
- Trait that implements behavior common to all collections based on the high-order function `foreach`.
- Provides an abstract method that generates an `Iterator`.

```
def foreach[U](f: Elem => U): Unit  
def iterator: Iterator[T]
```

- `Iterator` defines abstract methods `hasNext` and `next`.
- **Seq:** Provides a method `apply` for indexing (numerical index).
- **Buffer:** Used to create sequences by appending (`+=`), prepending (`+=:`), inserting, updating, or removing elements.
- **LinearSeq:** Efficient implementation of `head`, `tail`, `isEmpty`. Most other operations are linear.
- Many types provide an `apply` method in the companion object that should be used to create collection instances (default implementation will be used).

```
val seq = Seq(1,2,3)           // Vector  
val mutSeq = mutable.Seq(1, 3, 2) // ArrayBuffer
```

Parallel Collections



- The method `par` converts a sequential collection into its parallel counterpart.
- Methods of parallel collections like `map`, `foreach`, `filter`, `max`, `sum`, etc. execute in parallel.
- Example:

```
def seqSumSquares(seq: Seq[Int]) = seq.map(i => i*i).sum
def parSumSquares(seq: ParSeq[Int]) = seq.map(i => i*i).sum
val numbers = 1 to 1000000
val seqSum = seqSumSquares(numbers)
val parSum = parSumSquares(numbers.par)
```

Specifics of Parallel Collection

- Parallelizable collections

- Parallel *splitters* recursively split a collection into smaller parts:

```
trait Splitter[+T] extends Iterator[T]:  
  def split: Seq[Splitter[T]]
```

- The asymptotic running time of `split` must be $O(\log n)$. Otherwise, the recursive division would produce too much overhead.
- *Parallelizable collections*: Array, Vector, HashSet, HashMap, etc.
- *Non-parallelizable collections*: List, Stream

- Side effects in parallel operations

```
var sum = new AtomicLong(0)  
numbers.par foreach { i => sum += i }
```

- When multiple threads use variables, access must be synchronized, or atomic variables must be used.
- Parallel collection operation operations must be associative.